OwnFeed

Take the Algorithm Back

Business Plan — Confidential Draft — 2026

1. Executive Summary

OwnFeed is an open-source, AI-powered personal content curation tool that puts algorithmic control back in the hands of individuals. Instead of having their information diet shaped by engagement-maximizing platforms, users define their own curation logic in plain language — and the AI surfaces content accordingly from RSS, newsletters, podcasts, YouTube, Reddit, and the broader open web.

The product is intentionally open source. The soul of the tool must match its function: a platform that claims to give people control cannot itself become another dependency. Open source is both a principled commitment and a distribution strategy — the audience most motivated by algorithmic autonomy are the same people who evangelize and contribute to tools they trust.

Sustainability follows from community. The path to revenue is a hosted managed tier for individuals and an enterprise/B2B tier for teams, newsrooms, and organizations — a proven open core model. We build the open source version first, let the community validate the product, and monetize the convenience and scale layers second.

The core proposition in one sentence:
"Write your own algorithm in plain language. OwnFeed surfaces what you actually care about — not what keeps you scrolling."

2. Problem Statement

Every major content platform — YouTube, Instagram, TikTok, LinkedIn, Facebook — operates an algorithmic feed. Those algorithms share a common design principle: they optimize for engagement, which translates directly to ad revenue. The result is a systematic bias toward content that provokes strong reactions: outrage, anxiety, tribalism. Depth, nuance, and genuine learning are structurally disadvantaged.

Users are aware of this. Surveys consistently show widespread distrust of algorithmic feeds and a desire for more control. But awareness without an alternative produces learned helplessness. The tools that exist today fall short:

- RSS readers (Feedly, NetNewsWire) aggregate content from sources users already know, but have no intelligence layer and require manual curation of sources
- AI assistants (ChatGPT, Perplexity) answer questions but are not feed tools — they don't aggregate, rank, or surface content proactively
- Niche apps (Matter, Readwise Reader) serve long-form readers but don't address the broader feed problem or offer user-defined curation logic

The gap is clear: no tool currently lets a person define their own curation values in plain language and have an AI apply those values across a broad, multi-source content landscape. That is what OwnFeed does.

3. Market Opportunity

Target Audience

The initial audience is technical and privacy-conscious users — developers, researchers, knowledge workers, and founders who are already aware of algorithmic manipulation and motivated to address it. This is the Hacker News / self-hoster demographic: not large in absolute numbers, but highly influential, vocal, and likely to contribute, share, and evangelize.

The broader addressable market expands as the product matures: journalists, educators, researchers, policy professionals, and eventually any organization that wants to manage its team's information environment.

Why Now

- Public trust in algorithmic platforms is at a historic low following years of documented manipulation, mental health research, and regulatory scrutiny
- LLM capabilities have only recently reached the quality threshold where a natural language curation prompt can be reliably interpreted and acted upon
- Open protocols (Bluesky / AT Protocol, ActivityPub) are maturing and explicitly support custom algorithmic feeds as a first-class feature
- The post-Twitter fragmentation of social media has increased appetite for personal tools that aggregate across platforms rather than living inside one

Competitive Positioning

OwnFeed does not compete with platforms. It sits above them — an intelligence layer that aggregates open content and transforms it according to user-defined values. The competitive moat is not a proprietary dataset or exclusive API access. It is the quality of the personal preference model and the trust of an open source community.

4. Product Overview

Core Architecture

- **Content sources:** Aggregation layer

Pulls from RSS feeds, podcast directories, newsletter/email inboxes, YouTube public API, Reddit API, and open web sources. Designed for extensibility — new source connectors can be contributed by the community.

- **Curation engine:** Personal preference model

Users define their algorithm in plain language. Example: “I care about political economy, space, and emerging technology. Surface things that would genuinely surprise me. Deprioritize outrage. Nothing after 9pm that will spike my cortisol.” The LLM ranks and filters incoming content against this stated model.

- **Transparency:** Explainability layer

Every surfaced item shows why it was selected. Users can inspect, edit, and override their preference model at any time. No black box.

- **Shared taste models:** Community curation prompts

A library of user-contributed curation prompts — a researcher’s reading diet, a journalist’s source filter, a parent’s “news without doom” setup. Fork and remix any prompt as your starting point.

Technical Approach

- Open source core: self-hostable, MIT or Apache 2.0 license, GitHub-native development
- Local-first option: runs against local LLMs (Llama, Mistral, Phi) for fully private, zero-cost operation
- Bring-your-own-API-key: users who want cloud LLM quality connect their own OpenAI or Anthropic key
- Browser extension layer: transforms YouTube, Reddit, and X feeds in-browser without requiring platform API access
- Email wedge: dedicated inbox address for newsletter aggregation — zero platform dependency

5. Build Phases

Phase	Focus	Timeline	Key Deliverables
Phase 1 Foundation	Open source MVP on open web content	Months 1–3	• RSS + podcast aggregation • Plain-language curation prompt engine • Basic explainability (why surfaced)

			• GitHub repo, MIT license • Email newsletter wedge
Phase 2 Extension	Platform reach via browser extension	Months 4–6	• Chrome/Firefox extension • YouTube, Reddit, X feed transformation • User preference model persistence • Community prompt library v1
Phase 3 Data Ownership	User data import & richer taste models	Months 7–9	• Google Takeout import (YouTube, Gmail) • Twitter/X archive import • Taste model refinement from history • Bluesky / AT Protocol native feed
Phase 4 Monetization	Hosted tier & enterprise offering	Months 10–18	• Managed hosted tier (individuals) • Enterprise deployment package • Team & admin features • SLA, support, onboarding

6. Business Model

Open Core Structure

The product follows the proven open core model — a free, fully functional open source base with paid tiers for convenience and scale. The individual user is never the monetization target. The $0 tier is not a liability; it is the top of the funnel.

Open Source	Hosted Individual	Hosted Team	Enterprise / B2B
Free	**$5–10/mo**	**$15–30/mo per seat**	**Custom contract**
• Self-hosted • Full features • Bring own API key • Community support	• Managed hosting • No setup required • Cloud LLM included • Email support	• Team dashboard • Shared prompts • Admin controls • Priority support	• Custom deployment • SSO / security • SLA guarantee • Dedicated onboarding

Secondary Revenue Streams

- One-time packaged download ($5–15): pre-configured installer for non-technical users who want to self-host without setup friction

- GitHub Sponsors and open source grants (Mozilla Foundation, NLnet, Open Technology Fund)
- Dual licensing: MIT for individuals, commercial license for businesses building on top of OwnFeed

7. Milestones & Proposed Timeline

Milestone	Target Date	Success Criteria
MVP shipped to GitHub	Month 3	Working RSS + prompt curation engine, public repo, basic docs
First 100 GitHub stars	Month 4–5	Organic discovery via Hacker News, Product Hunt, or dev communities
First 500 active installs	Month 6	Measured via opt-in telemetry or GitHub release downloads
Browser extension launched	Month 6	Chrome Web Store listing, YouTube/Reddit feed transformation functional
Community prompt library v1	Month 6–7	20+ contributed curation prompts, fork/remix workflow live
Bluesky native feed	Month 9	OwnFeed listed as a custom feed algorithm on AT Protocol
Hosted tier launched	Month 12	Paying individual subscribers; infrastructure cost-neutral
First enterprise customer	Month 15–18	Newsroom, research org, or B2B team on custom contract

8. Funding Strategy

OwnFeed does not require external investment to reach a functional open source product. The initial build is feasible as a part-time effort alongside existing SIGNITEK consulting work, keeping burn low and preserving equity and mission integrity.

Phase 1–2: Bootstrapped

Development funded by existing consulting revenue. No external capital required. Estimated time investment: 8–12 hours per week across a founding team of 2–3 people.

Phase 3–4: Community & Grants

As the product finds its community, pursue:

- GitHub Sponsors: $500–2,000/month realistic at 500–1,000 active users with strong values alignment
- Open source infrastructure grants: Mozilla Foundation MOSS, NLnet NGI, Open Technology Fund — all have funded comparable privacy/autonomy tools
- Early hosted tier revenue: reinvested directly into infrastructure and development

Phase 4+: Selective Raise (Optional)

If enterprise traction materializes and growth requires dedicated engineering resources, a small seed raise ($250K–$500K) from mission-aligned investors becomes viable — but only if it does not compromise the open source commitment or community trust. This is a decision point, not a requirement.

> **Guiding principle on funding:**
> *The goal is a sustainable, independent product — not a venture-scale exit. Capital is a tool to serve the mission, not the mission itself. We raise only what we need to do the work.*

9. Future Action Items

Immediate (Next 30 Days)

- Finalize working product name and secure domain / GitHub organization
- Define MVP feature scope: minimum viable curation engine that demonstrates the value proposition
- Assess partner interest: determine level of involvement from Nick and Brian
- Choose technical stack: evaluate local-first (Node.js / Python) vs. cloud-native architecture for v1
- Register on Bluesky developer program and Space-Track equivalent for any open protocol integrations

Short-Term (Months 1–3)

- Build and ship MVP to GitHub with documentation and a clear README that explains the vision
- Submit to Hacker News ‘Show HN’ and relevant communities to gather early feedback
- Establish contribution guidelines and community norms early — open source culture is set in the first 90 days
- Identify 5–10 potential early users willing to test and give feedback

Medium-Term (Months 4–9)

- Iterate on curation quality based on user feedback — this is the core product differentiator
- Build browser extension and submit to Chrome Web Store
- Launch community prompt library and create onboarding flow around it
- Apply for first open source grant (Mozilla MOSS or NLnet)

- Evaluate Bluesky custom feed as a distribution and validation channel

Long-Term (Months 10+)

- Assess demand signals for hosted tier and launch if warranted
- Identify first enterprise or B2B prospects from existing networks
- Evaluate dual-license structure if commercial use cases emerge organically
- Reassess team structure and funding needs based on growth trajectory

---

A note on intent

OwnFeed is not a business-first project. It is a thing worth building first — and a business worth building second, on the right terms. The measure of success in Phase 1 is not revenue. It is whether the people who care most about this problem show up, use it, and make it better.

If that happens, everything else follows.